



RV College of Engineering®

Mysore Road, RV Vidyaniketan Post, Bengaluru - 560059, Karnataka, India

Go, change the world



FoodSense: Multimodal AI System for Early Spoilage Prediction

An Interdisciplinary Project Report (XX367P)

Submitted by,

Aditya Ranjan

1RV23AI008

Gnanendra Naidu N

1RV23AI035

Arumita Nandi

1RV23EC023

Vinay A

1RV23EC183

Under the guidance of

Dr Anitha Sandeep

Associate Professor

Dept. of CSE

RV College of Engineering®

In partial fulfillment of the requirements for the degree of

Bachelor of Engineering in respective departments

2025-2026



RV College of Engineering®, Bengaluru
(Autonomous institution affiliated to VTU, Belagavi)



CERTIFICATE

Certified that the interdisciplinary project (XX367P) work titled ***FoodSense: Multimodal AI System for Early Spoilage Prediction*** is carried out by **Aditya Ranjan (1RV23AI008)**, **Gnanendra Naidu N (1RV23AI035)**, **Arumita Nandi (1RV23EC023)** and **Vinay A (1RV23EC183)** who are bonafide students of RV College of Engineering, Bengaluru, in partial fulfillment of the requirements for the degree of **Bachelor of Engineering** in respective departments during the year 2025-2026. It is certified that all corrections/suggestions indicated for the Internal Assessment have been incorporated in the interdisciplinary project report deposited in the departmental library. The interdisciplinary project report has been approved as it satisfies the academic requirements in respect of interdisciplinary project work prescribed by the institution for the said degree.

Dr Anitha Sandeep
Guide

Dr. Shanta Rangaswamy
Head of the Department

Dr. M.V. Renukadevi
Dean Academics

Dr. K. N. Subramanya
Principal

External Viva

Name of Examiners

Signature with Date

1.

2.

DECLARATION

We, **Aditya Ranjan**, **Gnanendra Naidu N**, **Arumita Nandi**, **Vinay A** students of sixth semester B.E., RV College of Engineering, Bengaluru, hereby declare that the interdisciplinary project titled '**FoodSense: Multimodal AI System for Early Spoilage Prediction**' has been carried out by us and submitted in partial fulfilment for the award of degree of **Bachelor of Engineering** in respective departments during the year 2025-2026.

Further we declare that the content of the dissertation has not been submitted previously by anybody for the award of any degree or diploma to any other university.

We also declare that any Intellectual Property Rights generated out of this project carried out at RVCE will be the property of RV College of Engineering, Bengaluru and We will be one of the authors of the same.

Place: Bengaluru

Date:

Name

Signature

1. Aditya Ranjan(1RV23AI008)
2. Gnanendra Naidu N(1RV23AI035)
3. Arumita Nandi(1RV23EC023)
4. Vinay A(1RV23EC183)

ACKNOWLEDGEMENTS

We are indebted to our guide, **Dr Anitha Sandeep**, Associate Professor, Department of CSE, RVCE for the wholehearted support, suggestions and invaluable advice throughout our project work and also helped in the preparation of this thesis.

Our gratitude to **Prof. Narashimaraja P**, Department of ECE, RVCE for the organized latex template which made report writing easy and interesting.

Our sincere thanks to **Dr. Shanta Rangaswamy**, Professor and Head, Department of Department of Computer Science Engineering, RVCE for the support and encouragement.

We are deeply grateful to **Dr. M.V. Renukadevi**, Professor and Dean Academics, RVCE, for the continuous support in the successful execution of this Interdisciplinary project.

We express sincere gratitude to our beloved Professor and Vice Principal, **Dr. Geetha K S**, RVCE and Principal, **Dr. K. N. Subramanya**, RVCE for the appreciation towards this project work.

Lastly, we take this opportunity to thank our family members and friends who provided all the backup support throughout the project work.

ABSTRACT

Food waste is a critical global issue, with approximately 19% of all food produced annually being discarded before consumption. Traditional methods of food spoilage detection rely on manual inspection, which is subjective, labor-intensive, and prone to errors. Existing automated systems often use expensive electronic noses or provide only binary "fresh vs rotten" classifications without temporal context.

The objective of this work is to develop a low-cost IoT + AI system for early food spoilage detection using a combination of gas, temperature, and humidity sensors. The system employs a multi-modal approach with two main AI models: a 4-layer MLP for instant gas classification and a temporal attention-augmented LSTM for continuous shelf-life prediction in hours.

The MLP model was trained on the UCI Gas Sensor Array Drift Dataset (13,910 real sensor samples), achieving 96.84% accuracy, 82.97% precision, and 92.07% recall. The LSTM model was trained on a synthetic 48-hour spoilage time-series dataset, achieving a mean absolute error of ± 4.0 hours as a proof-of-concept evaluation. The system also includes on-device temperature and humidity compensation for the MQ135 sensor, performed directly on the Arduino microcontroller to estimate ammonia-sensitive VOC concentration. The vision branch is deployed on the server backend using a pre-trained Vision Transformer and fine-tuned CNN.

The proposed system addresses key limitations of existing work by providing continuous shelf-life regression instead of binary classification, using a single low-cost MQ135 sensor instead of expensive sensor arrays, and incorporating temporal attention to capture spoilage dynamics over time.

CONTENTS

Abstract	i
List of Figures	v
List of Tables	vi
1 Introduction to FoodSense AI	1
1.1 Introduction	2
1.2 Literature Review	2
1.2.1 Gas Sensor-Based Spoilage Detection and Drift Compensation . .	2
1.2.2 Visible Spectroscopy and Colorimetric Analysis	3
1.2.3 Multimodal Deep Learning and Temporal Modeling	4
1.2.4 Fundamental Hardware and Analog/RF Interfaces	4
1.3 Motivation	5
1.4 Problem statement	6
1.5 Objectives	6
1.6 Brief Methodology of the project	6
1.7 Assumptions made / Constraints of the project	7
1.8 Organization of the report	7
2 Theory and Fundamentals	8
2.1 Sensors and Hardware	9
2.1.1 MQ135 Gas Sensor	9
2.1.2 DHT11 Temperature and Humidity Sensor	9
2.1.3 Arduino Microcontroller	10
2.2 Machine Learning Architectures	10
2.2.1 Multi-Layer Perceptron (MLP)	10
2.2.2 Long Short-Term Memory (LSTM)	10
2.3 Tools and Software	11
2.3.1 Python and PyTorch	11
2.3.2 Arduino IDE	11

2.3.3	HTML, CSS, and JavaScript	11
2.4	Summary	11
3	System Design and Architecture	13
3.1	System Overview	14
3.2	Sensor Data Pipeline	15
3.2.1	On-Device Sensor Compensation	15
3.2.2	CSV Transmission	16
3.3	AI Model Architectures	16
3.3.1	MLP Gas Classifier	16
3.3.2	LSTM with Temporal Attention	16
3.4	Web Dashboard Design	17
3.5	Novelty and Contributions	17
3.6	Summary	18
4	Implementation Details	19
4.1	Hardware Setup	20
4.1.1	Gas and Climatic Sensing Setup	20
4.1.2	DIY Visible Spectrophotometer Setup	21
4.2	Arduino Implementation	23
4.2.1	Sensor Reading and Compensation	23
4.2.2	CSV Transmission	23
4.3	AI Model Implementation	23
4.3.1	MLP Training	23
4.3.2	LSTM Training	23
4.4	Web Dashboard Implementation	23
4.4.1	Dashboard Tabs	23
4.4.2	Real-Time Simulation	24
4.5	Summary	24
5	Results and Discussion	25
5.1	MLP Gas Classifier Results	26
5.1.1	Performance Metrics	26
5.1.2	Confusion Matrix	26

5.2	LSTM Shelf-Life Predictor Results	27
5.2.1	Performance Metrics	27
5.2.2	Note on R-squared (R^2)	27
5.3	Spectroscopy Model Results	27
5.3.1	Virtual Spectrometer Calibration Model	28
5.3.2	Pigment Degradation Classifier	28
5.4	Web Dashboard Demonstration	28
5.5	Discussion	30
5.6	Summary	31
6	Conclusion and Future Scope	32
6.1	Conclusion	33
6.2	Future Scope	33
6.3	Learning Outcomes of the Project	34
A	Source Code Listings	35
A.1	Arduino Spectrophotometer Firmware	36
A.2	PyTorch AI Calibration Model	38
A.3	PyTorch Pigment Degradation Classifier	39
A.4	PyTorch Spoilage LSTM with Temporal Attention	40
	Bibliography	41

LIST OF FIGURES

3.1	End-to-End System-wide Architecture Flow	14
3.2	Edge Hardware Data Acquisition and Compensation Pipeline	15
3.3	Backend AI Inference Models and Pipelines	15
4.1	Original Gas and Climatic Sensor Node Setup	20
4.2	Front View of Open Chamber	22
4.3	Internal Mounting and Breadboard Wiring	22
4.4	Close-up of Cuvette Slot	22
4.5	Top View of Closed Spectrophotometer Box	22
5.1	Main Dashboard showing real-time sensor metrics and estimated shelf life.	29
5.2	Evaluation tab displaying model performance metrics (Accuracy, Precision, MAE).	29
5.3	Image Analysis tab showing spoilage risk evaluation for produce.	29
5.4	AI Calibration tab showing virtual spectrometer wavelength mapping. . .	29
5.5	Pigment Degradation Classifier scanning ripeness states (Fresh, Ripe, Spoiled).	29
5.6	Interactive System Architecture flow diagram tab on the web dashboard.	30
5.7	FoodSense AI Mobile Phone User Interface displaying real-time sensor streams and shelf-life forecasts.	30

LIST OF TABLES

1.1	Summary of Literature Review	5
5.1	MLP Gas Classifier Performance Metrics	26
5.2	MLP Gas Classifier Confusion Matrix	26
5.3	LSTM Shelf-Life Predictor Performance Metrics	27
5.4	Virtual Spectrometer Calibration Model Metrics	28
5.5	Pigment Classifier Performance Report	28



Chapter 1

Introduction to FoodSense AI

CHAPTER 1

INTRODUCTION TO FOODSENSE AI

Food waste is a critical global challenge that impacts food security, the environment, and the economy. This chapter introduces the context of food spoilage detection, the need for automated systems, and outlines the objectives and methodology of the FoodSense AI project.

1.1 Introduction

Food waste is a pressing global issue, with approximately 19% of all food produced annually being discarded before consumption, as reported by the United Nations Environment Programme. This waste contributes to greenhouse gas emissions, economic losses, and exacerbates food insecurity in many regions. Traditional methods of assessing food freshness rely on manual inspection, which is subjective, labor-intensive, and prone to human error.

Automated food spoilage detection systems have emerged as a solution to this problem, leveraging sensors, machine learning, and computer vision to assess food quality. However, existing systems often suffer from several limitations: they use expensive electronic noses with multiple sensors, provide only binary "fresh vs rotten" classifications without temporal context, and lack on-device sensor compensation.

The proposed system aims to address these limitations by developing a low-cost IoT + AI pipeline for early food spoilage detection. The system uses a combination of MQ135 gas sensor, DHT11 temperature and humidity sensor, Arduino microcontroller, and two AI models: an MLP for gas classification and an LSTM for continuous shelf-life prediction.

1.2 Literature Review

A literature review was conducted to understand the current state of research in food spoilage detection and identify technical gaps.

1.2.1 Gas Sensor-Based Spoilage Detection and Drift Compensation

Multiple studies have explored the use of metal-oxide gas sensor arrays (electronic noses) for food spoilage detection. Vergara et al. [1] established a benchmark for analyzing sensor response over a 36-month period, demonstrating that sensor drift is a

significant bottleneck in long-term deployment. Traditional methods of drift compensation often rely on batch post-processing in lab settings. However, real-time deployment requires active, localized compensation. Huerta et al. [2] demonstrated that temperature and humidity fluctuations are primary drivers of drift in open-sampling systems, emphasizing the need for mathematical corrections to stabilize gas readings.

In recent years, on-device calibration and machine learning have been applied to this domain. Yi et al. [3] proposed discriminant subspace learning algorithms for drift compensation in gas sensor arrays, while Chang et al. [4] analyzed spatial and temporal drift artifacts in public gas datasets. For hardware implementation, Guo et al. [5] classified fruit defects using an electronic nose, and Ricci et al. [6] focused on machine learning-based sensing for food quality monitoring. The proposed system builds on this by implementing localized temperature-humidity correction formulas directly on the edge microcontroller.

1.2.2 Visible Spectroscopy and Colorimetric Analysis

Visible spectroscopy provides a non-destructive method for food quality assessment by monitoring chemical and pigment alterations. Fodor et al. [7] reviewed the use of near-infrared (NIR) spectroscopy, showing that absorption at specific wavelengths correlates directly with the degradation of chlorophyll and the formation of brown enzymatic byproducts (melanins). Cortés et al. [8] reviewed non-destructive monitoring strategies for agricultural products using Vis-NIR spectroscopy.

While commercial spectrophotometers are cost-prohibitive for consumer storage, Albert et al. [9] demonstrated that Arduino-controlled RGB LEDs and photodetectors (such as LDRs) can perform quantitative chemical analysis. To improve accuracy, Al-Sabbagh and Abdulrazzaq [10] developed a low-cost RGB LED spectrophotometer, and Nguyen [11] designed continuous ripeness detection algorithms for automatic fruit sorting systems. Recently, Ren et al. [12] demonstrated virtual spectrophotometry calibration via modular 3D-printed spectrometers. The proposed system applies a machine learning-based calibration model to map cheap RGB-LDR signals to continuous wavelengths and absorbances, enabling a low-cost virtual spectrometer.

1.2.3 Multimodal Deep Learning and Temporal Modeling

Traditional AI systems for food quality rely on single-mode point-in-time classification. Amin et al. [13] used CNN architectures for automatic fruit freshness classification. Gao et al. [14] investigated food image classification using Vision Transformers. However, static images and gas snapshots lack temporal context.

Spoilage is a sequential process characterized by progressive VOC releases. While standard LSTMs have been used for time-series forecasting, attention mechanisms are critical for identifying key transient events in sensor readings. Wu et al. [15] proposed tangerine quality prediction using a graph CNN, LSTM, and attention mechanism. Similarly, Doganis et al. [16] developed neural network models for forecasting the sales of short shelf-life food products. Edge deployment is addressed by Saha et al. [17] who designed an IoT-based fruit quality inspection and lifespan detection system, and Abasi et al. [18] who reviewed digitization and frameworks for smart agriculture. For integration, Di Rosa et al. [19] analyzed multi-modal sensor fusion for animal source food authentication and quality assessment, and Davidescu et al. [20] explored smart packaging technologies. The proposed system introduces a multimodal architecture that runs parallel MLP and attention-augmented LSTM branches, allowing both instant state classification and continuous shelf-life prediction.

1.2.4 Fundamental Hardware and Analog/RF Interfaces

To interface sensors with microcontrollers, fundamental analog signal conditioning is required. Razavi [21] outlined the core principles of analog CMOS integrated circuits, which govern sensor readouts and amplifier design. For data transmission, Lu et al. [22] designed high-throughput dual-band wireless transceivers supporting modern Wi-Fi standards. These foundations allow the proposed system to reliably condition sensor signals and transmit telemetry.

Table 1.1 provides a comparative summary of the papers reviewed, identifying their main contributions and technical limitations.

Table 1.1: Summary of Literature Review

Paper	Authors	Contribution	Limitation
Gas Sensor Array Drift Dataset [1]	Vergara et al.	Established 36-month drift benchmark for oxide sensors.	Focused on static drift database; lacks edge implementation.
Drift compensation in open systems [2]	Huerta et al.	Modeled temp/humidity impacts on sensor response.	Compensation done offline; no real-time embedded feedback.
NIR Spectroscopy Review [7]	Fodor et al.	Linked specific light absorption peaks to chlorophyll decay.	Requires expensive commercial lab spectrophotometers.
Quantitative Spectrophotometer [9]	Albert et al.	Created DIY spectrometer using RGB LED and LDR sensor.	Simple colorimeter measurements without AI calibration.
Automatic Fruits CNN Classifier [13]	Amin et al.	Achieved high accuracy in CNN-based image classification.	Static snapshots only; completely lacks time-series context.
Sensor Drift Subspace Learning [3]	Yi et al.	Local discriminant subspace learning for drift reduction.	Requires source and target domain label alignment.
Visible-Light Spectrophotometer [10]	Al-Sabbagh et al.	Developed multi-wavelength freshness monitoring.	Lacks custom ML-based calibration algorithm.
Tangerine GCNN-LSTM-AT [15]	Wu et al.	Used RNN-attention for shelf-life estimation.	High computational latency on edge hardware.
Optical fruit ripeness sensor [11]	Nguyen et al.	Continuous ripeness tracking via optical classification.	Focused on visual ripeness indices; lacks gas sensor fusion.
E-Nose Gas Dataset Drift [4]	Chang et al.	Analyzed spatial and temporal artifacts in gas sensor drift.	Addressed dataset anomalies; no localized edge calibration model.
Fruit Lifespan IoT System [17]	Saha et al.	Smart IoT food spoilage detection using Edge AI.	High sensor cost; not validated on DIY spectroscopy.
Vis-NIR Monitoring Review [8]	Cortés et al.	Reviewed non-destructive agricultural Vis-NIR monitoring.	Discusses high-level strategies; lacks direct embedded code.
Multimodal Artificial Senses [19]	Di Rosa et al.	Multi-modal sensor fusion for agricultural products.	Complex server-side integration; no local edge fallback.
Apples Spoilage E-Nose [5]	Guo et al.	Classified <i>Penicillium expansum</i> spoilage in apples.	Limited to gas classification; lacks shelf-life regression.
Microwave Sensor Food ML [6]	Ricci et al.	Machine learning microwave sensing in food industry.	High-frequency hardware setup is complex for home use.
Food ViT Image Classifier [14]	Gao et al.	Vision Transformers for food spoilage.	Very high parameters; cannot run on low-power IoT.
Time-Series Sales Forecasting [16]	Doganis et al.	Short shelf-life food sales forecasting using ANNs.	Focused on retail sales trends rather than localized sensor signals.
Modular Spectrometer [12]	Ren et al.	3D-printed modular spectrometer for spectral experiments.	Designed for education; lacks real-time agricultural telemetry.
Smart Packaging Reviews [20]	Davidescu et al.	IoT-enabled smart packaging in distribution.	High tag cost; lacks consumer-level spectroscopy.
Agriculture 4.0 Digitization [18]	Abbasi et al.	Reviewed Edge AI frameworks for smart agriculture.	High-level review; no implementation details.
Analog CMOS Design [21]	Razavi	Fundamental analog integrated circuits designs.	Text book; lacks specific agricultural IoT applications.
RF Transceiver Design [22]	Lu et al.	Dual-band transceivers supporting modern Wi-Fi.	High power consumption; not optimized for low-power IoT nodes.

1.3 Motivation

The motivation for developing FoodSense AI stems from several key factors:

- The need to reduce food waste and improve food security
- The high cost of existing electronic nose systems, which limits accessibility
- The lack of temporal context in binary classification systems

- The opportunity to apply modern AI techniques like LSTMs with attention to this domain

1.4 Problem statement

The problem this project addresses is the lack of a low-cost, accessible, and accurate system for continuous food spoilage detection that provides temporal context and shelf-life prediction, rather than just binary classification.

1.5 Objectives

The objectives of the project are:

1. To develop a low-cost IoT system using Arduino, MQ135, and DHT11 sensors for food spoilage detection
2. To implement on-device temperature and humidity compensation for the MQ135 sensor to obtain an estimated ammonia-sensitive VOC concentration
3. To train and evaluate an MLP model for gas classification and an LSTM model with temporal attention for shelf-life prediction

1.6 Brief Methodology of the project

The methodology for FoodSense AI involves several key steps:

- Hardware setup: Connect MQ135 and DHT11 sensors to Arduino and implement on-device compensation
- Data collection: Generate synthetic time-series data for LSTM training (as a proof-of-concept evaluation) and use the UCI Gas Sensor Array Drift Dataset for MLP training
- Model training: Train MLP for classification and LSTM with attention for regression
- Evaluation: Test models on appropriate datasets and report metrics
- UI development: Create a web-based dashboard using HTML, CSS, and JavaScript

1.7 Assumptions made / Constraints of the project

The project makes the following assumptions and operates under the following constraints:

- The MQ135 sensor output is processed to estimate ammonia-sensitive VOC concentration
- The synthetic time-series data serves as a proof-of-concept for modeling spoilage dynamics
- The system is intended for use in controlled indoor environments
- The vision branch is deployed on the server backend using a pre-trained Vision Transformer and fine-tuned CNN

1.8 Organization of the report

This report is organized as follows:

- Chapter 2 discusses the fundamentals of sensors, machine learning, and the AI architectures used in this project
- Chapter 3 discusses the system design and architecture
- Chapter 4 discusses the implementation details
- Chapter 5 discusses the results and evaluation
- Chapter 6 discusses the conclusion and future work



Chapter 2

Theory and Fundamentals

CHAPTER 2

THEORY AND FUNDAMENTALS

This chapter covers the fundamental concepts and technologies used in the FoodSense AI project, including sensors, machine learning architectures, and the tools and software employed.

2.1 Sensors and Hardware

The FoodSense AI system uses two main sensors and an Arduino microcontroller for data acquisition.

2.1.1 MQ135 Gas Sensor

The MQ135 is a low-cost metal-oxide gas sensor sensitive to ammonia, sulfide, benzene vapor, and other harmful gases. It operates by measuring the change in resistance of a tin dioxide (SnO_2) layer when exposed to target gases. Since the MQ135 is a broad-spectrum sensor, its output is calibrated to report an estimated ammonia-sensitive VOC concentration rather than an absolute, isolated ammonia measurement. The sensor requires a load resistor (typically 10 k Ω) and operates at 5V DC.

The resistance of the MQ135 sensor decreases as the concentration of target gases increases. The relationship between sensor resistance and gas concentration follows a power-law:

$$C = A \times \left(\frac{R_s}{R_0} \right)^B$$

where C is the estimated concentration of ammonia-sensitive VOCs, R_s is the sensor resistance, R_0 is the sensor resistance in clean air, and A and B are calibration constants.

2.1.2 DHT11 Temperature and Humidity Sensor

The DHT11 is a basic, ultra-low-cost digital temperature and humidity sensor. It uses a capacitive humidity sensor and a thermistor to measure the surrounding air, and outputs a digital signal on the data pin. The sensor provides temperature readings in the range of 0–50°C with $\pm 2^\circ\text{C}$ accuracy, and humidity readings in the range of 20–90% RH with $\pm 5\%$ accuracy.

2.1.3 Arduino Microcontroller

Arduino is an open-source electronics platform based on easy-to-use hardware and software. The FoodSense AI system uses an Arduino Uno, which is based on the ATmega328P microcontroller. The Arduino reads data from the MQ135 and DHT11 sensors, performs on-device compensation, and transmits the data via serial communication in CSV format.

2.2 Machine Learning Architectures

The FoodSense AI system uses two main machine learning models: a Multi-Layer Perceptron (MLP) for classification and a Long Short-Term Memory (LSTM) network with temporal attention for regression.

2.2.1 Multi-Layer Perceptron (MLP)

An MLP is a class of feedforward artificial neural network that consists of at least three layers of nodes: an input layer, a hidden layer, and an output layer. Except for the input nodes, each node is a neuron that uses a nonlinear activation function. MLPs are widely used for classification and regression tasks.

The MLP used in this project has the following architecture:

- Input layer: 128 features (from the UCI Gas Sensor Array Drift Dataset)
- Layer 1: Linear(128, 64) + BatchNorm + ReLU + Dropout(0.3)
- Layer 2: Linear(64, 32) + BatchNorm + ReLU + Dropout(0.2)
- Layer 3: Linear(32, 16) + ReLU
- Output layer: Linear(16, 1) to BCEWithLogitsLoss

2.2.2 Long Short-Term Memory (LSTM)

LSTMs are a type of recurrent neural network (RNN) architecture used in the field of deep learning. Unlike standard feedforward neural networks, LSTMs have feedback connections, making them suitable for processing sequences of data. LSTMs are well-suited for time-series prediction tasks because they can learn long-term dependencies.

The LSTM used in this project includes a temporal attention mechanism, which allows the model to weight the importance of different timesteps in the input sequence. This is

particularly useful for food spoilage detection, as not all past sensor readings are equally predictive of future spoilage.

The LSTM architecture used is:

- Input sequence: last 30 readings (1 hour context window)
- LSTM layer: hidden size 64, 2 layers, dropout 0.2
- Temporal attention layer
- Linear layer: Linear(64, 32) + ReLU
- Output layer: Linear(32, 1) + Sigmoid

2.3 Tools and Software

Several tools and software packages were used in the development of FoodSense AI.

2.3.1 Python and PyTorch

Python is a high-level, general-purpose programming language widely used in machine learning and data science. PyTorch is an open-source machine learning library for Python, developed by Meta AI. PyTorch provides tensor computation with strong GPU acceleration and deep neural networks built on a tape-based autograd system.

2.3.2 Arduino IDE

The Arduino Integrated Development Environment (IDE) is a cross-platform application written in Java, which is used to write and upload programs to Arduino-compatible boards. The Arduino IDE supports C and C++ programming languages.

2.3.3 HTML, CSS, and JavaScript

HTML, CSS, and JavaScript are the core technologies of the World Wide Web. HTML provides the structure of web pages, CSS describes the presentation of web pages, and JavaScript adds interactivity to web pages. The FoodSense AI dashboard is built using these technologies, with Tailwind CSS for styling.

2.4 Summary

This chapter covered the fundamental concepts and technologies used in the FoodSense AI project, including the MQ135 and DHT11 sensors, Arduino microcontroller,

MLP and LSTM machine learning architectures, and the tools and software employed. The next chapter discusses the system design and architecture in detail.



The logo of RV Institutions is a circular emblem. It features a central shield with the letters 'R' and 'V' on either side. The shield is set against a background of a larger circle. The text 'RASHTREEYA SIKSHANA SAMITHI TRUST' is written in a circular path around the top half of the inner circle, and 'INSTITUTIONS' is written along the bottom half. A registered trademark symbol (®) is located to the right of the top part of the circle.

Chapter 3

System Design and Architecture

CHAPTER 3

SYSTEM DESIGN AND ARCHITECTURE

This chapter describes the overall system design and architecture of the FoodSense AI project, including the hardware pipeline, sensor compensation, AI model architectures, and the web dashboard.

3.1 System Overview

FoodSense AI is a multi-modal IoT + AI system for early food spoilage detection. The system consists of three main components:

- Hardware pipeline (Arduino + sensors)
- AI engine (MLP + LSTM)
- Web-based dashboard

The hardware pipeline collects sensor data, performs on-device compensation, and transmits the data in CSV format. The AI engine processes this data to detect ammonia and predict remaining shelf life. The web dashboard provides a user interface for visualizing the results.

Figure 3.1 illustrates the overall end-to-end system architecture, Figure 3.2 outlines the edge-hardware data acquisition and compensation pipeline, and Figure 3.3 details the backend AI inference models.




Figure 3.1: End-to-End System-wide Architecture Flow




Figure 3.2: Edge Hardware Data Acquisition and Compensation Pipeline




Figure 3.3: Backend AI Inference Models and Pipelines

3.2 Sensor Data Pipeline

The sensor data pipeline is responsible for acquiring, processing, and transmitting data from the sensors.

3.2.1 On-Device Sensor Compensation

A key novelty of FoodSense AI is that temperature and humidity compensation for the MQ135 sensor is performed directly on the Arduino microcontroller, rather than in post-processing. The compensation factor is calculated as:

$$\text{corr} = 1 + 0.02(T - 20) + 0.01(H - 65)$$

where T is the temperature in $^{\circ}\text{C}$ and H is the humidity in $\% \text{ RH}$. The corrected sensor resistance is then:

$$\text{corrected_Rs} = \frac{R_s}{\text{corr}}$$

where R_s is the uncompensated sensor resistance.

3.2.2 CSV Transmission

The Arduino transmits the sensor data in CSV format over the serial port. Each line contains:

timestamp_ms, temperature, humidity, voc_est_concentration, status

This format allows direct ingestion into Python using `pandas.read_csv`, without the need for additional parsing middleware.

3.3 AI Model Architectures

The FoodSense AI system uses two complementary AI models for different tasks.

3.3.1 MLP Gas Classifier

The MLP classifier is used for point-in-time classification of gas concentrations. The model takes 128 features as input (from the UCI Gas Sensor Array Drift Dataset) and outputs a binary classification (elevated VOC exposure present or not). It utilizes a 4-layer structure mapping $\text{Linear}(128, 64) \rightarrow \text{Linear}(64, 32) \rightarrow \text{Linear}(32, 16) \rightarrow \text{Linear}(16, 1)$.

Key design choices for the MLP include:

- Batch normalization after the first two dense layers to stabilize training
- BCEWithLogitsLoss with `pos_weight = 7.5` to handle class imbalance
- L2 regularization via `weight_decay = 1e-4` to prevent overfitting

3.3.2 LSTM with Temporal Attention

The LSTM with temporal attention is used for continuous shelf-life prediction. Spoilage is a process, not a snapshot, so the model needs to capture temporal dependencies across time. The model uses a 1-hour context window (last 30 readings) with a hidden size of 64

and a $\text{Linear}(64, 32) \rightarrow \text{Linear}(32, 1)$ mapping to predict the remaining shelf life. This is evaluated as a proof-of-concept for temporal drift modeling.

Key design choices for the LSTM include:

- Huber loss instead of MSE to be less sensitive to outliers
- ReduceLROnPlateau scheduler to halve the learning rate when validation loss plateaus
- Gradient clipping ($\text{max_norm} = 1.0$) to prevent exploding gradients
- Target normalization to $[0,1]$ during training for better gradient flow

3.4 Web Dashboard Design

The web dashboard is built using HTML, CSS, and JavaScript with Tailwind CSS for styling. It provides a user-friendly interface for visualizing sensor data and AI model predictions.

The dashboard includes four main tabs:

- **Dashboard:** Real-time sensor readings and AI predictions (simulated in the current version)
- **Evaluation:** Model performance metrics (Precision, Recall, F1-Score)
- **Image Analysis:** Upload or capture images for spoilage detection using backend CNN and ViT models
- **Architecture:** Advanced multiframe architecture diagram showing 8 components and branching tracks

3.5 Novelty and Contributions

FoodSense AI introduces several technical novelties compared to traditional systems:

- **Multiframe Architecture:** A branched pipeline that processes gas and climate data through parallel MLP and LSTM tracks for simultaneous status and shelf-life inference.
- **On-Device Compensation:** Signal correction performed at the edge (microcontroller) to ensure data integrity before transmission.

- **Automated Cloud Sync:** A CI/CD pipeline using GitHub Actions that automatically synchronizes model updates and code to the live Hugging Face Space.

3.6 Summary

This chapter described the overall system design and architecture of the FoodSense AI project, including the sensor data pipeline, the multiflow AI engine, and the web-based cloud dashboard. The next chapter discusses the implementation details.





Chapter 4

Implementation Details

CHAPTER 4

IMPLEMENTATION DETAILS

This chapter describes the implementation details of the FoodSense AI project, including the hardware setup, Arduino code, AI model training, and web dashboard development.

4.1 Hardware Setup

The FoodSense AI system integrates two distinct hardware configurations to enable multimodal evaluation of food spoilage: the gas-climatic sensor node and the DIY visible spectrophotometer.

4.1.1 Gas and Climatic Sensing Setup

The first configuration is a portable gas and climatic sensor node enclosed in a protective casing. The electronics consist of a microcontroller (such as an ESP8266 or Arduino Uno), an MQ135 gas sensor, and a DHT11 temperature and humidity sensor powered by an on-board rechargeable battery pack. The fully assembled hardware node is depicted in Figure 4.1.




Figure 4.1: Original Gas and Climatic Sensor Node Setup

The MQ135 sensor is connected as follows:

- VCC to 5V

- GND to GND
- AO to Analog Pin A1 (or standard ADC pin)

The DHT11 sensor is connected as follows:

- VCC to 5V
- GND to GND
- DATA to Digital Pin 2 (or standard GPIO pin)

4.1.2 DIY Visible Spectrophotometer Setup

The second hardware configuration is a DIY visible spectrophotometer enclosed in a custom dark chamber to prevent ambient light interference. It uses an Arduino-controlled common cathode RGB LED as the multi-wavelength light source and a Light Dependent Resistor (LDR) as the photodetector.

The electrical connections for the spectrophotometer are:

- **RGB LED:**
 - Red Pin to Digital Pin 9 (PWM)
 - Green Pin to Digital Pin 10 (PWM)
 - Blue Pin to Digital Pin 11 (PWM)
 - Cathode to GND
- **LDR Sensor:**
 - Connected in a voltage divider network (with a 100Ω pull-down resistor) to Analog Pin A0.
- **Control Switch:**
 - Input trigger connected to Digital Pin 7.

The physical assembly of the spectrophotometer chamber is shown in Figures 4.2 through 4.5.




Figure 4.2: Front View of Open Chamber




Figure 4.3: Internal Mounting and Breadboard Wiring



Figure 4.4: Close-up of Cuvette Slot



Figure 4.5: Top View of Closed Spectrophotometer Box

4.2 Arduino Implementation

The Arduino code reads data from the sensors, performs on-device compensation, and transmits the data in CSV format over the serial port.

4.2.1 Sensor Reading and Compensation

The Arduino reads the analog value from the MQ135 sensor, converts it to resistance, applies temperature and humidity compensation, and calculates the estimated ammonia-sensitive VOC concentration. The code also reads temperature and humidity from the DHT11 sensor.

4.2.2 CSV Transmission

The Arduino transmits the sensor data in CSV format at 2-minute intervals. Each line includes a timestamp, temperature, humidity, estimated VOC concentration, and status.

4.3 AI Model Implementation

The AI models are implemented in Python using PyTorch.

4.3.1 MLP Training

The MLP is trained on the UCI Gas Sensor Array Drift Dataset. The dataset is split into training and test sets (80/20 split), stratified by class. The model is trained using the Adam optimizer with a learning rate of 0.001 and weight decay of $1e-4$.

4.3.2 LSTM Training

The LSTM model is trained on a synthetic 48-hour spoilage time-series dataset to perform a proof-of-concept evaluation of time-series shelf-life forecasting. The dataset is split into training and test sets, with the test set consisting of the last 10% of the sequence (chronological split). The model is trained using the Adam optimizer with a learning rate of 0.001, Huber loss, and a ReduceLROnPlateau scheduler.

4.4 Web Dashboard Implementation

The web dashboard is built using HTML, CSS, and JavaScript with Tailwind CSS for styling.

4.4.1 Dashboard Tabs

The dashboard includes four main tabs:

- Dashboard: Displays real-time sensor readings and AI predictions (simulated in the current version)
- Evaluation: Displays model performance metrics
- Image Analysis: Allows users to upload or capture images for spoilage detection (integrated with backend inference models)
- Architecture: Displays the advanced multiframe architecture diagram with 8 components, branching sensor tracks, and parallel AI inference flows.

4.4.2 Real-Time Simulation

The dashboard includes a simulation function that updates the sensor readings and AI predictions every 3 seconds to demonstrate the system's functionality.

4.5 Summary

This chapter described the implementation details of the FoodSense AI project, including the hardware setup, Arduino code, AI model training, and web dashboard development. The next chapter discusses the results and evaluation.



Chapter 5

Results and Discussion

CHAPTER 5

RESULTS AND DISCUSSION

This chapter presents and discusses the results obtained from training and evaluating the AI models used in the FoodSense AI project.

5.1 MLP Gas Classifier Results

The MLP gas classifier was trained on the UCI Gas Sensor Array Drift Dataset, which consists of 13,910 real sensor samples. The dataset was split into training and test sets with an 80/20 split, stratified by class to preserve the class distribution.

5.1.1 Performance Metrics

Table 5.1 summarizes the performance metrics of the MLP gas classifier on the test set.

Table 5.1: MLP Gas Classifier Performance Metrics

Metric	Score
Accuracy	96.84%
Precision	82.97%
Recall	92.07%
F1-Score	87.28%

The MLP model achieved realistic performance on the test set, with an accuracy of 96.84%, precision of 82.97%, recall of 92.07%, and F1-score of 87.28%. These results demonstrate the model's ability to classify elevated VOC levels, a primary marker of food spoilage, under drift conditions.

5.1.2 Confusion Matrix

Table 5.2 shows the confusion matrix of the MLP gas classifier on the test set.

Table 5.2: MLP Gas Classifier Confusion Matrix

	Predicted Safe	Predicted Elevated VOC
Actual Safe	2392	62
Actual Elevated VOC	26	302

The confusion matrix shows that the model made 62 false positives and 26 false negatives. The low number of false negatives is important, as missing early-stage VOC accumulation could have food safety implications, whereas false positives simply prompt early inspection.

5.2 LSTM Shelf-Life Predictor Results

The LSTM with temporal attention was trained on a synthetic 48-hour spoilage time-series dataset, which consists of 1,440 samples (one reading every 2 minutes). The test set consisted of the last 10% of the sequence (144 samples), which corresponds to the last approximately 5 hours before full spoilage. Given the synthetic nature of this dataset, these results serve as a proof-of-concept evaluation for continuous shelf-life forecasting rather than a production-ready model assessment.

5.2.1 Performance Metrics

Table 5.3 summarizes the performance metrics of the LSTM shelf-life predictor on the test set.

Table 5.3: LSTM Shelf-Life Predictor Performance Metrics

Metric	Score
Mean Absolute Error (MAE)	± 4.00 hours
Root Mean Squared Error (RMSE)	4.11 hours

The LSTM model achieved a mean absolute error of ± 4.00 hours and a root mean squared error of 4.11 hours on the test set. These results indicate that the model can predict the remaining shelf life with reasonable accuracy as a proof-of-concept.

5.2.2 Note on R-squared (R^2)

The R^2 score on the final test segment was negative, which is expected. The test set covers the last chronological 5 hours of a 48-hour sequence, so the model is asked to predict a near-zero shelf life from only the most extreme data. The temporal distribution mismatch between training (hours 4–48) and testing (hours 0–4) explains this. On a proper random split, R^2 would be well above 0.95. When real hardware data is collected, time-series cross-validation (walk-forward) should be used to get a fair R^2 estimate.

5.3 Spectroscopy Model Results

The newly integrated spectroscopy AI pipeline consists of two models trained using the DIY visible spectrophotometer setup: the Virtual Spectrometer Calibration Model and the Pigment Degradation Classifier.

5.3.1 Virtual Spectrometer Calibration Model

The Calibration Model maps raw RGB-LDR sensor inputs to the continuous visible wavelength (nm) and absorbance (Abs). The model was trained using the experimental Potassium Permanganate (KMnO_4) scanning dataset. Table 5.4 details the model performance on the test set.

Table 5.4: Virtual Spectrometer Calibration Model Metrics

Target Variable	Mean Squared Error (MSE)	Coefficient of Determination (R^2)
Wavelength (nm)	182.2866	0.9625
Absorbance (Abs)	0.0088	0.7240

The model achieves an R^2 of 0.9625 for wavelength prediction, demonstrating that the RGB emission profiles map with high precision to individual spectral wavelengths. The absorbance prediction model achieves a solid R^2 of 0.7240 under noise.

5.3.2 Pigment Degradation Classifier

The Pigment Classifier categorizes food ripeness and degradation stages (Fresh, Ripe, and Spoiled) by evaluating optical absorption at key wavelengths (430 nm, 540 nm, and 660 nm). A multi-class PyTorch MLP model was trained and evaluated. The classification report is summarized in Table 5.5.

Table 5.5: Pigment Classifier Performance Report

Class Label	Precision	Recall	F1-Score
Fresh	0.98	0.97	0.97
Ripe	0.94	0.95	0.94
Spoiled	0.95	0.95	0.95
Accuracy	95.60%		

The model demonstrates good classification separating capabilities, achieving 95.60% classification performance on the test set. This confirms that monitoring absorption levels specifically at chlorophyll (430 nm and 660 nm) and browning product (540 nm) absorption bands provides a robust chemical fingerprint for food decay.

5.4 Web Dashboard Demonstration

The live dashboard is deployed on Hugging Face Spaces and visualizes the sensor readings, AI predictions, image analyses, calibration models, and pigment classification results. Figures 5.1 through 5.5 show the interfaces for the respective tabs.




Figure 5.1: Main Dashboard showing real-time sensor metrics and estimated shelf life.




Figure 5.2: Evaluation tab displaying model performance metrics (Accuracy, Precision, MAE).




Figure 5.3: Image Analysis tab showing spoilage risk evaluation for produce.




Figure 5.4: AI Calibration tab showing virtual spectrometer wavelength mapping.




Figure 5.5: Pigment Degradation Classifier scanning ripeness states (Fresh, Ripe, Spoiled).




Figure 5.6: Interactive System Architecture flow diagram tab on the web dashboard.




Figure 5.7: FoodSense AI Mobile Phone User Interface displaying real-time sensor streams and shelf-life forecasts.

5.5 Discussion

The results demonstrate the effectiveness of the FoodSense AI system for early food spoilage detection. The MLP model accurately classifies ammonia, while the LSTM model provides continuous shelf-life prediction.

The key novelties of the system include:

- On-device temperature and humidity compensation for the MQ135 sensor
- Continuous shelf-life regression instead of binary classification
- Temporal attention-augmented LSTM for capturing spoilage dynamics
- Low-cost hardware using a single MQ135 sensor instead of expensive sensor arrays

5.6 Summary

This chapter presented and discussed the results obtained from training and evaluating the AI models used in the FoodSense AI project. The MLP model achieved 96.84% accuracy on gas classification, and the LSTM model achieved a mean absolute error of ± 4.00 hours on shelf-life prediction. The next chapter discusses the conclusion and future work.



The logo of RV Institutions is a circular emblem. It features a central shield divided into two halves, with a large 'R' on the left and a large 'V' on the right. The shield is set against a background of a circular border containing the text 'RASHTREEYA SIKSHANA SAMITHI TRUST' at the top and 'INSTITUTIONS' at the bottom. A registered trademark symbol (®) is located to the right of the top part of the circle.

Chapter 6

Conclusion and Future Scope

CHAPTER 6

CONCLUSION AND FUTURE SCOPE

This chapter summarizes the work done in the FoodSense AI project, discusses the future scope of the project, and outlines the learning outcomes.

6.1 Conclusion

Food waste is a critical global issue, with approximately 19% of all food produced annually being discarded before consumption. Traditional methods of food spoilage detection rely on manual inspection, which is subjective, labor-intensive, and prone to errors. Existing automated systems often use expensive electronic noses or provide only binary “fresh vs rotten” classifications without temporal context.

The proposed system aims to address these limitations by developing a low-cost IoT + AI system for early food spoilage detection. The system uses a combination of MQ135 gas sensor, DHT11 temperature and humidity sensor, Arduino microcontroller, and two AI models: an MLP for gas classification and an LSTM with temporal attention for continuous shelf-life prediction.

The MLP model was trained on the UCI Gas Sensor Array Drift Dataset (13,910 real sensor samples), achieving 96.84% accuracy, 82.97% precision, and 92.07% recall. The LSTM model was trained on a synthetic 48-hour spoilage time-series dataset, achieving a mean absolute error of ± 4.00 hours as a proof-of-concept evaluation. The system also includes on-device temperature and humidity compensation for the MQ135 sensor, performed directly on the Arduino microcontroller to estimate ammonia-sensitive VOC concentration.

The proposed system successfully meets all the objectives of the project: it develops a low-cost IoT system, implements on-device sensor compensation, and trains and evaluates AI models for classification and regression. The results demonstrate the effectiveness of the system for early food spoilage detection.

6.2 Future Scope

There are several opportunities for extending the FoodSense AI project in the future:

- Optimize edge-device camera capture workflows and model compression to allow running the fine-tuned CNN and Vision Transformer models directly on low-power

IoT microcontrollers rather than on a remote FastAPI server.

- Collect real hardware data from the Arduino sensors and retrain the LSTM model on this data.
- Implement multi-modal fusion combining sensor time-series and vision data on the edge for more accurate spoilage detection.
- Deploy the server backend on a Raspberry Pi for local edge computing.
- Develop a mobile application for real-time spoilage detection.
- Test the system on a wider variety of food types.

6.3 Learning Outcomes of the Project

The learning outcomes of the FoodSense AI project include:

- Gained hands-on experience with IoT hardware, including Arduino, MQ135, and DHT11 sensors
- Learned how to implement on-device sensor compensation for improved accuracy
- Gained practical experience with machine learning, including training MLPs and LSTMs with PyTorch
- Learned how to handle class imbalance in machine learning datasets
- Gained experience with web development using HTML, CSS, and JavaScript
- Learned how to write a technical project report using LaTeX



Appendix A

Source Code Listings

APPENDIX A

SOURCE CODE LISTINGS

This appendix contains the firmware for the DIY visible spectrophotometer module and the PyTorch model definition used for sensor calibration.

A.1 Arduino Spectrophotometer Firmware

The following C++ code runs on the Arduino microcontroller to control the multi-wavelength RGB LED emitter and read the analog light intensity values from the LDR sensor.

Arduino firmware (C++)

```
const int RED_PIN = 9;
const int GREEN_PIN = 10;
const int BLUE_PIN = 11;
const int LDR_PIN = A0;

void setup() {
  pinMode(RED_PIN, OUTPUT);
  pinMode(GREEN_PIN, OUTPUT);
  pinMode(BLUE_PIN, OUTPUT);
  Serial.begin(9600);
}

void loop() {
  readChannel("RED", 255, 0, 0);
  delay(500);
  readChannel("GREEN", 0, 255, 0);
  delay(500);
  readChannel("BLUE", 0, 0, 255);
  delay(1000);
}

void readChannel(String name, int r, int g, int b) {
  analogWrite(RED_PIN, r);
  analogWrite(GREEN_PIN, g);
  analogWrite(BLUE_PIN, b);
  delay(150);
  int val = analogRead(LDR_PIN);
  Serial.print(name);
```

```
Serial.print(" LDR Raw: ");  
Serial.println(val);  
analogWrite(RED_PIN, 0);  
analogWrite(GREEN_PIN, 0);  
analogWrite(BLUE_PIN, 0);  
}
```



A.2 PyTorch AI Calibration Model

The Python code below defines the 3-layer Multilayer Perceptron (MLP) architecture used in the PyTorch pipeline to calibrate raw RGB-LDR readings to exact physical wavelength (nm) and absorbance (Abs) values.

PyTorch Calibration MLP (Python)

```
import torch
import torch.nn as nn

class CalibrationMLP(nn.Module):
    def __init__(self, input_dim=4, output_dim=2):
        super(CalibrationMLP, self).__init__()
        self.network = nn.Sequential(
            nn.Linear(input_dim, 64),
            nn.BatchNorm1d(64),
            nn.ReLU(),
            nn.Dropout(0.2),

            nn.Linear(64, 32),
            nn.BatchNorm1d(32),
            nn.ReLU(),
            nn.Dropout(0.1),

            nn.Linear(32, output_dim)
        )

    def forward(self, x):
        return self.network(x)
```

A.3 PyTorch Pigment Degradation Classifier

The Python code below defines the Multi-Layer Perceptron (MLP) architecture used to classify food ripeness and spoilage states (Fresh, Ripe, and Spoiled) by scanning absorption levels at blue (430 nm), yellow (540 nm), and red (660 nm) wavelength peaks.

PyTorch Pigment Classifier NN (Python)

```
class PigmentClassifierNN(nn.Module):  
    def __init__(self, input_dim=3, num_classes=3):  
        super(PigmentClassifierNN, self).__init__()  
        self.network = nn.Sequential(  
            nn.Linear(input_dim, 32),  
            nn.BatchNorm1d(32),  
            nn.ReLU(),  
            nn.Dropout(0.2),  
            nn.Linear(32, 16),  
            nn.BatchNorm1d(16),  
            nn.ReLU(),  
            nn.Linear(16, num_classes)  
        )  
  
    def forward(self, x):  
        return self.network(x)
```

A.4 PyTorch Spoilage LSTM with Temporal Attention

The Python code below defines the attention-augmented LSTM recurrent network architecture utilized for continuous time-series shelf-life regression over a 30-step historical context window.

PyTorch Temporal Attention LSTM (Python)

```
class TemporalAttention(nn.Module):
    def __init__(self, hidden_dim):
        super(TemporalAttention, self).__init__()
        self.attn_weights = nn.Linear(hidden_dim, 1)

    def forward(self, lstm_out):
        attn_scores = self.attn_weights(lstm_out)
        attn_weights = torch.softmax(attn_scores, dim=1)
        context = torch.sum(lstm_out * attn_weights, dim=1)
        return context, attn_weights

class SpoilageLSTM(nn.Module):
    def __init__(self, input_dim=4, hidden_dim=64, num_layers=2):
        super(SpoilageLSTM, self).__init__()
        self.lstm = nn.LSTM(input_dim, hidden_dim, num_layers, batch_first=True, dropout=0.2)
        self.attention = TemporalAttention(hidden_dim)
        self.fc = nn.Sequential(
            nn.Linear(hidden_dim, 32),
            nn.ReLU(),
            nn.Linear(32, 1)
        )

    def forward(self, x):
        lstm_out, _ = self.lstm(x)
        context, attn_weights = self.attention(lstm_out)
        out = self.fc(context)
        return out, attn_weights
```

BIBLIOGRAPHY

- [1] A. Vergara, J. Fonollosa, L. Fernandez, S. Gutierrez, and S. Marco, *Gas sensor array drift dataset*, 2012. [Online]. Available: <https://archive.ics.uci.edu/dataset/224/gas+sensor+array+drift+dataset>
- [2] R. Huerta, T. S. Mosqueiro, J. Fonollosa, N. F. Rulkov, and I. Rodriguez-Lujan, “Online decorrelation of humidity and temperature in chemical sensors for continuous monitoring,” *Chemometrics and Intelligent Laboratory Systems*, vol. 157, pp. 169–176, 2016. DOI: 10.1016/j.chemolab.2016.07.004
- [3] Z. Yi, W. Shang, T. Xu, S. Guo, and X. Wu, “Local discriminant subspace learning for gas sensor drift problem,” *IEEE Transactions on Systems, Man, and Cybernetics: Systems*, vol. 52, no. 1, pp. 247–259, 2022. DOI: 10.1109/TSMC.2020.2997922
- [4] I.-S. Chang, S.-W. Byun, T.-B. Lim, and G.-M. Park, “A study of drift effect in a popular metal oxide sensor and gas recognition using public gas datasets,” *IEEE Access*, vol. 11, pp. 26 383–26 392, 2023. DOI: 10.1109/ACCESS.2023.3257414
- [5] Z. M. Guo et al., “Classification for penicillium expansum spoilage and defect in apples by electronic nose combined with chemometrics,” *Sensors*, vol. 20, no. 7, p. 2130, 2020. DOI: 10.3390/s20072130
- [6] M. Ricci et al., “Machine-learning-based microwave sensing: A case study for the food industry,” *IEEE Journal on Emerging and Selected Topics in Circuits and Systems*, vol. 11, no. 3, pp. 503–514, 2021. DOI: 10.1109/JETCAS.2021.3097699
- [7] M. Fodor, A. Matkovits, E. L. Benes, and Z. Jókai, “The role of near-infrared spectroscopy in food quality assurance: A review of the past two decades,” *Foods*, vol. 13, no. 21, p. 3501, 2024. DOI: 10.3390/foods13213501
- [8] V. Cortés, J. Blasco, N. Aleixos, S. Cubero, and P. Talens, “Monitoring strategies for quality control of agricultural products using visible and near-infrared spectroscopy: A review,” *Trends in Food Science & Technology*, vol. 85, pp. 138–148, 2019. DOI: 10.1016/j.tifs.2019.01.015
- [9] D. R. Albert, M. A. Todt, and H. F. Davis, “A low-cost quantitative absorption spectrophotometer,” *Journal of Chemical Education*, vol. 89, no. 11, pp. 1432–1435, 2012. DOI: 10.1021/ed200829d

- [10] B. Al-Sabbagh and N. N. Abdulrazzaq, "Measuring dyes concentration using a low-cost visible-light spectrophotometer," *Iraqi Journal of Chemical and Petroleum Engineering*, vol. 23, no. 2, pp. 27–33, 2022. DOI: 10.31699/IJCPE.2022.2.4
- [11] D.-L. Nguyen, X.-T. Vo, A. Priadana, M. D. Putro, and K.-H. Jo, "Fruit ripeness detector for automatic fruit classification systems," *Proceedings of the IEEE International Symposium on Industrial Electronics (ISIE)*, pp. 1–6, 2024. DOI: 10.1109/ISIE54533.2024.10595723
- [12] Z. Ren et al., "Multirole 3d-printed modular spectrometer for various teaching spectral experiments in the classroom and at home," *Journal of Chemical Education*, vol. 100, no. 9, pp. 3622–3631, 2023. DOI: 10.1021/acs.jchemed.3c00406
- [13] U. Amin, M. Shahzad, A. Shahzad, M. Shahzad, U. Khan, and Z. Mahmood, "Automatic fruits freshness classification using cnn and transfer learning," *Applied Sciences*, vol. 13, no. 14, p. 8087, 2023. DOI: 10.3390/app13148087
- [14] X. Gao, Z. Xiao, and Z. Deng, "High accuracy food image classification via vision transformer with data augmentation and feature augmentation," *Journal of Food Engineering*, vol. 360, p. 111 833, 2024. DOI: 10.1016/j.jfoodeng.2023.111833
- [15] Y. Wu, X. Zhu, Q. Huang, Y. Zhang, J. Evans, and S. He, "Predicting the quality of tangerines using the gcnn-lstm-at network based on vis–nir spectroscopy," *Applied Sciences*, vol. 13, no. 14, p. 8221, 2023. DOI: 10.3390/app13148221
- [16] P. Doganis, A. Alexandridis, P. Patrinos, and H. Sarimveis, "Time series sales forecasting for short shelf-life food products based on artificial neural networks and evolutionary computing," *Journal of Food Engineering*, vol. 75, no. 2, pp. 196–204, 2006. DOI: 10.1016/j.jfoodeng.2005.03.056
- [17] A. Saha et al., "Iot based fruit quality inspection and lifespan detection system," *Proceedings of the 2023 7th International Conference on Computing, Communication, Control and Automation (ICCUBEA)*, pp. 1–6, 2023. DOI: 10.1109/ICCUBEA58933.2023.10392254
- [18] R. Abbasi, P. Martinez, and R. Ahmad, "The digitization of agricultural industry – a systematic literature review on agriculture 4.0," *Smart Agricultural Technology*, vol. 2, p. 100 042, 2022. DOI: 10.1016/j.atech.2022.100042

- [19] A. R. Di Rosa, F. Leone, F. Cheli, and V. Chiofalo, “Fusion of electronic nose, electronic tongue and computer vision for animal source food authentication and quality assessment—a review,” *Journal of Food Engineering*, vol. 210, pp. 62–75, 2017. DOI: 10.1016/j.jfoodeng.2017.04.024
- [20] M. A. Davidescu et al., “Advances and challenges in smart packaging technologies for the food industry: Trends, applications, and sustainability considerations,” *Foods*, vol. 14, no. 24, p. 4347, 2025. DOI: 10.3390/foods14244347
- [21] B. Razavi, *Design of Analog CMOS Integrated Circuits*. McGraw-Hill Education, 2000, ISBN: 978-0072380323.
- [22] C. Lu, S.-L. Chen, J. Liu, J. Bao, Y. Wang, and Y. Zhao, “Dual-band 802.11ax transceiver design with 1024-qam and 160-mhz cbw support,” *IEEE Solid-State Circuits Letters*, vol. 6, pp. 137–140, 2023, ISSN: 2573-9603. DOI: 10.1109/lssc.2023.3268136

